Writeup for third project of
CMSC 420: "Data Structures"
Section 0101, Fall 2019

Theme: Hash Tables

On-time deadline: Sunday, 11-03, 11:59pm (midnight)
Late deadline (30% penalty): Wednesday, 11-06,
11:59pm (midnight)

Contents

1 Overview

In this project, you will have to implement an abstraction over a *phonebook*; A collection of pairs of type $< Full\_Name, Phone\_Number >$. Your phonebook will support **both** name-based search **and** phone-based search. See Figure 1 for a pictorial view of the project.
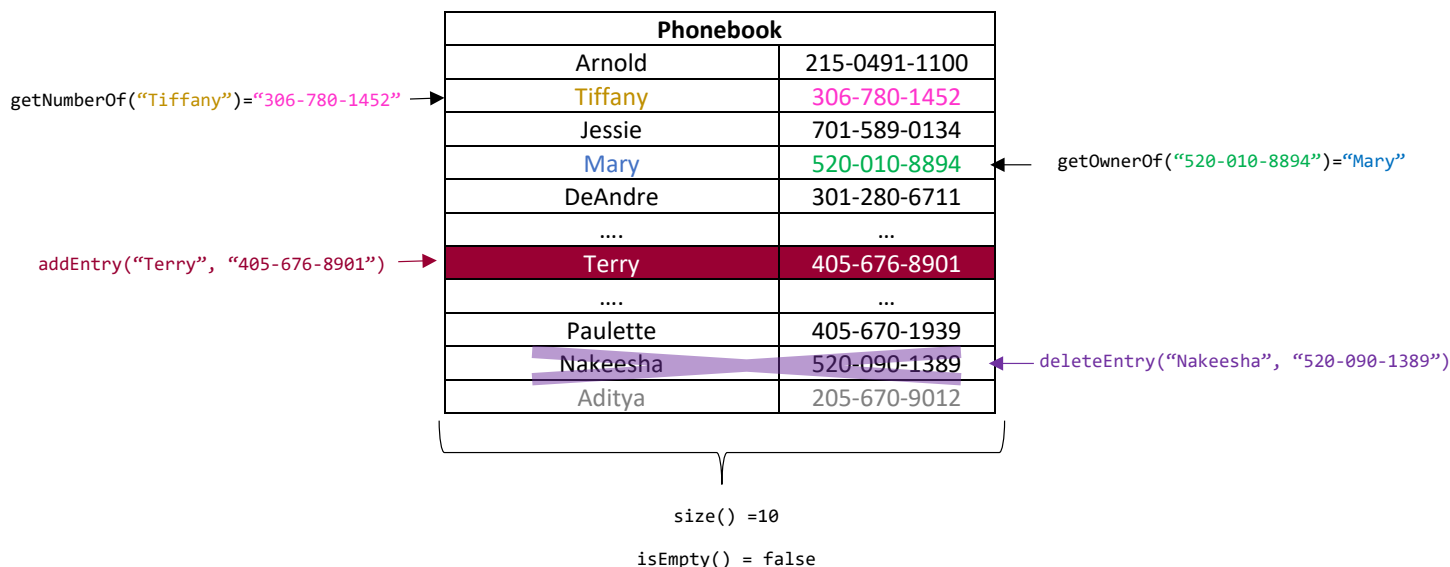

Figure 1: A high-level view of how your phonebook is supposed to work.

To make both types of searches efficient, your phonebook will internally maintain **a pair of hash tables** from `String`s to `String`s: One will have the **person's name as a key** and the **phone number as a value**, and the other one will have the **phone number as a value** and the **name as a key**! In your simple phonebook, **entry uniqueness** is **guaranteed**: Every person has **exactly** one phone number, and every phone number is associated with **exactly one** person.

2 Getting Started

You should first pull the starter code from our GitHub repo. To do this, **before** you pull, fill up your staging area with a `git add -A`, commit with a `git commit -m YOUR_COMMIT_MSG` and run a `git pull` to pull the new project's code. If there are any kinds of merge conflicts (unlikely) you can solve them manually by looking at the affected source code files and erasing the lines you don't want to change. After that, you should study the JavaDocs and source code of `Phonebook` to understand how your methods can be used (and, therefore, tested!). The classes you have to implement are under the package `hashes`.
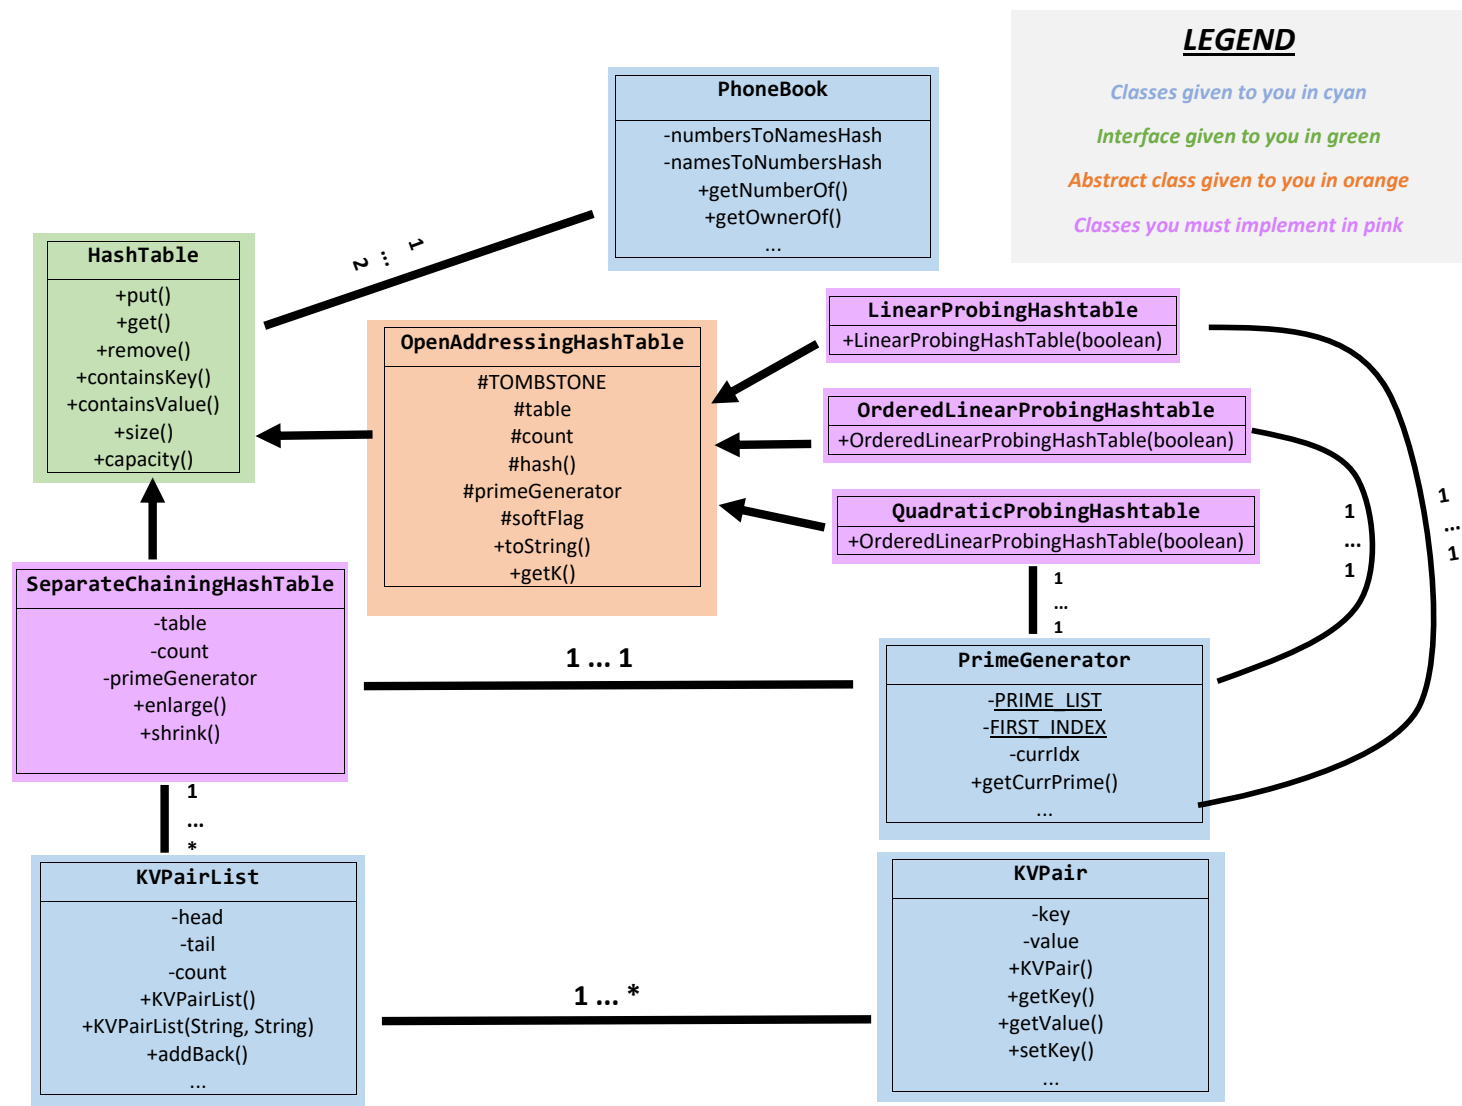
3 Provided code

3.1 Class hierarchy

Figure 2: A UML diagram describing the code structure of project 3. Simple lines reflect one-to-many (1 - *) "has - a" relations, while arrows show "is-a" relations (derived class, implemented interface, etc).

Figure 2 depicts the code structure for the project. The top-level class of the project is `Phonebook`. What is interesting about `Phonebook` is that **it has been implemented for you**! However, the methods of `Phonebook` depend on methods of the interface `HashTable`, which is implemented (directly or through `OpenAddressingHashTable` by the four classes which you will have to implement. What *you* will need to do is complete the implementation of these four classes so that their methods can support the methods of `Phonebook`. The Release Tests **primarily** test methods of `Phonebook` (approx. 90% of their code), while

3

a smaller number of tests check if you are implementing basic hash table functionality correctly (e.g resizings, see below).

The various methods of `Phonebook` will have to run in *amortized constant time* (except for `size()` and `isEmpty()`, which should run in *constant* time). This does *not* take into account the case of an insertion or deletion that results in a *resizing* of the array; we want amortized constant time *assuming* that a resizing does **not** occur in that particular operation. We **will** be checking your source code after submission to make sure you are **not** implementing the methods **inefficiently** (e.g logarithmic complexity, linear complexity, or even worse)!

In practice, the only way you can do this is by **not** consulting the hash function **at all** for your operations; just **looping over the entire table** until you either find the element (`remove`, `containsKey`) or you find an empty position (`put`). While this would indeed allow you to pass the tests, we will be **checking your submission** to make sure you **consult the hash function**! Implementing all operations as mentioned above would constrain them to be *linear time*, which is **unacceptable** for both the project (i.e no credit for this project) and Computer Science **as a whole**. That is correct; you will have disappointed the Science **as a whole**. That said, **not all** of the methods you implement make use of the hash function, which means that their complexity parameters will necessarily be different (can you find any such methods?).

You should fill in the `public` methods of `SeparateChainingHashTable`, `LinearProbing HashTable`, `OrderedLinearProbingHashTable` and `QuadraticProbingHashTable`. For the last three classes, you might find that some of the methods have the **exact same source code**. You would then benefit by making them `protected` methods in the `OpenAddressing HashTable` class, which those three classes extend!

3.2   Interfaces, `abstract` classes and the `protected` access modifier

In the code base, you might notice that `HashTable` is a Java *interface*. On the other hand, `OpenAddressingHashTable` is an `abstract` class. Abstract classes in Java are almost like interfaces, except they are allowed to contain fields and their members do **not** default to `public`. Similarly to interfaces, one **cannot** instantiate an `abstract` class.

The choice of making `OpenAddressingHashTable` an `abstract` class is deliberate; several methods and fields of **all** your Openly Addressed hash tables are common across **all** of these classes. Therefore, it makes sense to package them into *one place* and debug them in *one place*. Unfortunately, Java interfaces do **not** allow for storing fields, but only methods, which are also implicitly `public`. Essentially, the `HashTable` interface tells us what kinds of methods any `HashTable` instance **ought** to satisfy. For example, every `HashTable` instance needs to provide a method called `put`, with two `String` arguments, `key` and `value`, as well as a return value of type `Probes`. Refer to section 3.4 for a short diatribe on this small class. It should also answer questions of key containment, (`containsKey(String key)` and queries of its current stored key count (`size()`) and capacity (`capacity()`).

On the other hand, any Openly Addressed Hash Table needs to have some common fields and functionality. They all need an array over `KVPair` instances. They all need to answer queries of size and capacity in $\mathcal{O}(1)$. They can all benefit from an overriding of `toString()`, which we provide for you and is very useful for debugging.[1] Therefore, this entire piece of common functionality can (and should) be packaged in **one common place**, and this place in our code base is `OpenAddressingHashTable`.

You might notice that all the fields and methods of `OpenAddressingHashTable` besides `toString()` are labeled with the `protected` access modifier. `protected` essentially means: "visible by derived classes". In more detail, every *identifier* (name) of a field or a method that has been declared `protected` in a base class can be *straightforwardly* accessed from a derived class by its name, without any prepending of base class name or of any other name. This is very useful for your code! You might notice that `LinearProbingHashTable`, `OrderedLinearProbingHashTable` and `QuadraticProbingHashTable`, all classes that you **must** implement, have **no private fields** (but if you would like to add some, please go right ahead). In fact, our own implementation of the project does not add any extra fields in the classes (but we **do** use `private` methods for readability).

This is, of course, **not** a perfect approach towards building this code base. For example, you might notice that `SeparateChainingHashTable`, the only collision resolution method you have to implement that is not an Open Addressing method, *necessarily* has to declare some `private` fields that we also see in `OpenAddressingHashTable`. Also, somebody can argue that since `OrderedLinearProbingHashTable` has so many common characteristics (section 4 analyzes this in some detail) with `LinearProbingHashTable`, one could make the former of a subclass of the latter. You can come up with many different approaches of refactoring this code base, but **PLEASE DON'T, OTHERWISE YOU MIGHT END UP NOT PASSING ANY OF OUR UNIT TESTS!** The **only** thing you can do is add your own `private` fields or methods in the classes you have to implement and, in the case of the three Open Addressing methods you have to implement, you can move some of the common code you build as `protected` methods in `OpenAddressingHashTable`. See the comments at the very end of that class' definition for a relevant prompt.

The **tl;dr** of what you *can* and *cannot* change in the code base is this: unit tests test `public` functionality and they also need to be aware of type information at **compile-time**, since Java is a *strongly* typed language. *Are you in any way breaking the* `public` *methods' signatures and / or return types? Are you in any way altering the code base's hierarchy by making classes extending other classes and interfaces?* If the answer to **both** of these questions is **no**, you are good, otherwise you are **not** good. :)

3.3   Classes under `hashes`

Besides the classes you have to implement, you are given the following classes under the package `hashes`:

---

[1]So, if you are *persistent* about **not** using the debugger, **at least print your table** before coming to office hours, please ;(

- `CollisionResolver`: A simple `enum` which only contains four named fields, disambiguating between the various collision resolution methods that you will have to implement.

- `HashTable`: The top-level interface discussed in sections 3.1 and 3.2.

- `OpenAddressingHashTable`: The abstract class discussed in sections 3.1 and 3.2.

3.4 Classes under `utils`

The package `utils` will be *indispensable* to you. Here is a short description of what every class in the package does. Refer to the Javadocs for a *complete* and *concrete* description of arguments, return values, `Exception`s thrown, etc. Without consulting the Javadocs, you are **extremely likely to not be passing several tests**. For example, some of our tests expect that you will `throw` particular `Exception` instances in certain cases: the JavaDoc is your **only** guide in those situations! This list is just a **high-level** understanding of the methods.

- `KVPair`: An important abstraction for **Key-Value pairs**. Time and again in this class we have conveyed to you that we are assuming that the *value* with which a particular (and unique) *key* is associated is what we are *really* interested in, and the keys are only useful for somehow organizing the potentially infinite set of values, such that we can insert, search and delete as efficiently as possible, taking into considerations issues of cache locality, where our memory resides, how hard these K-V stores are to implement, etc.

  `KVPair` implements exactly that: it is a simple class which encapsulates **both** the key **and** the value into one place so that we can access the value from the key in $\mathcal{O}(1)$. In C/C++, we would probably have replaced it with a `struct`.

- `KVPairList`: An explicitly coded linked list that holds `KVPair` instances. It is only useful for `SeparateChainingHashTable`. If you are wondering why we opted for coding an entirely new list instead of simply using one of `Java`'s several generic `List`s, refer to section 6 for an explanation of how `Java` treats arrays of generic types, such as `KVPairList`. The short answer is: **not well at all**.

- `KVPairListTests`: A simple unit testing library for `KVPairList`.

- `PrimeGenerator`: A *very* important **singleton** class which controls the re-sizing parameters for **all** of your `HashTable` instances. In class, we have discussed the importance of keeping the size of your hash table as a *prime* number. This class helps us with that. In particular, you should study the JavaDocs for `getNextPrime()` and `getPreviousPrime()`, since you will certainly be using those methods for your own purposes. Both of these methods run in *constant* time, since we have already stored

a large list of primes as a `static` shared field of the `PrimeGenerator` class, and the various primes can be accessed by indexing into that field.

- `PrimeGeneratorTests`: A simple unit testing library for `PrimeGenerator`.

- `NoMorePrimesException`: A type of `RuntimeException` that `PrimeGenerator` uses when it runs out of primes to provide to an application.

- `Probes`: Arguably **the most important class for your testing and understanding**. The most important operations of `HashTable` instances, which are `put`, `get` and `remove`, **all** return `Probes` instances. These instances contain the *value* of the key that was inserted, sought or deleted (`null` in case of a failure of any kind), and, crucially, the **number of probes** that it took for the operation to succeed **or fail**. In this way, we can determine if you have understood how the collision resolution methods are supposed to work! Specifically, what the **length** of a collision chain ought to be *dynamically*, during execution of the code with successive operations on the same `HashTable` instance. A reminder that even an *immediately successful or unsuccessful insertion, deletion or search* **still** counts as one probe.

4 Collision Resolution methods

Given that the number of keys to store (e.g individual ATM transactions over the entire state of Maryland for a large bank organization) is **enormous** and the available space to store them in computer memory is *much* smaller, collisions are much, even with an excellent hash function. It therefore becomes important to develop *collision resolution methods*, whose job is to determine how an insertion of a key that *collides* with an existing key is resolved.

4.1 Separate Chaining

The most natural collision resolution method that we examine is **Separate Chaining**. In your code, this collision resolution method corresponds to `SeparateChainingHashTable`. An example of this method can be seen in Figure 3. Note that it is not **necessary** that we employ a linked list, or any list for that matter, for implementing the collision "chains". We could just as well use an AVL Tree, a Red-Black or B-Tree or a SkipList! The benefit of using a linked list for our collision chains is that we can insert very fast (by adding to the front or, in this project, by adding to the back with a `tail` pointer). The drawback is that we have linear time for search, but with M relatively large and a good hash function, we are hoping that the collision chains will, on average, have length $\frac{n}{M}$, which is still linear time but offers the favorable constant of $\frac{1}{M}$.

Figure 3: An example of the Separate Chaining collision resolution technique with integer keys. We see that the insertion of 140 collides with that of 7, while the fourth cell (at index 3) has had three collisions total to take care of.

As seen in Figure 2, `SeparateChainingHashTable` is the **only** class you have to implement which is **not** derived from `OpenAddressingHashTable`. This is intuitive: this method is the only one that stores the keys outside the table. It is wasteful in terms of memory, though, since for a table of capacity c we are spending $4c$ bytes (for 32-bit `Java` references). If $c = 1,000,000,000$, that is $4GB$ used just to *point* to the data that interests us! However, it is *very* easy to implement, it is *very* useful for estimating the quality of our hash function and it is also very useful if we want to retrieve a pointer to a different container as our value (e.g AVL Tree, a linked list, another hash...)

4.2   Linear Probing

Linear Probing (hereafter referred to as **LP**) is the oldest and simplest Open Addressing collision resolution method. It is a well-studied technique with some very attractive properties, first introduced and analyzed by Donald Knuth in 1963. An example of some insertions into a table that employs LP to store some **integers** is shown in 4. The hash function employed is a simple "modular" hash: $h(i) = i \bmod M$.


Figure 4: An example of the Linear Probing collision resolution technique in a hash table that stores integer keys.

8

Every time a collision is encountered, the algorithm keeps going forward into the table, wrapping around when required, to find an appropriate place to insert the new key into. 19, 4 and 16 are inserted collision-free, paying the minimum of only one probe, but 6, 176 and 1714 will be inserted only after enduring two, three and four probes respectively! Also note that this is the **maximum** number of insertions the hash table can accommodate before resizing; the next insertion is **guaranteed** to trigger a resizing of the table according to its resizing policy.

We will now offer a mathematical formalization of how LP works. Suppose that our hash function is $h(k)$, where k is some input key. Let also $i \geq 1$ be an integer that denotes the i^{th} probe that we have had to endure during our search for an empty cell in the table. We select $i \geq 1$ because, remember, the minimum #probes is 1 (one) , even for an **unsuccessful search**! Assuming that our hash table employs LP, the following **memory allocation function** $m_{lp}(k, i)$, returns the *actual cell index* of the i^{th} probe:

$$m_{lp}(k, i) = (h(k) + (i - 1)) \bmod M$$

This means that LP will probe the following memory addresses in the original hash table:

$$h(k) \bmod M, (h(k) + 1) \bmod M, (h(k) + 2) \bmod M, (h(k) + 3) \bmod M, \ldots$$

which fits intuition. For example, in the hash table shown in Figure 5, if we wanted to insert the key 22, we would have the **sequential** memory allocations: $m_{lp}(22, 1) = h(22) + (1 - 1) \bmod 11 = 22 \bmod 11 = 0$, $m_{lp}(22, 2) = \cdots = 1$ and $m_{lp}(22, 3) = 3$.


Figure 5: Examples of various memory allocations for two integer keys to be inserted into an hash table that employs Linear Probing.

On the other hand, if we wanted to insert the key 9, we would only need the single allocation $m_{lp}(9,1) = 9$, since cell 9 is empty. Of course, we could also compute $m_{lp}(9,2) = 10$ or $m_{lp}(9,3) = 0$, but there is no reason to, since m_{lp} gave us an empty address in the first probe.

LP has been praised for its **simplicity**, **excellent cache locality** and **theoretical properties** when employing a quality hash function (we will discuss those in lecture). But what would happen if we were to employ a relatively *poor* hash function?

To demonstrate what can happen, let's envision the following scenario. We have the following simple hash function for lowercase English characters:

$$h_{char}(c) = (int)c - 97$$

Since lowercase 'a' has the decimal value 97 in the ASCII table, we can subtract 97 to "zero-index" our hash functions for lowercase English characters. Table 1 can provide you with a reference of English characters throughout the rest of this writeup.

Character (a-m)	a	b	c	d	e	f	g	h	i	j	k	l	m
Value of h$_{char}$	0	1	2	3	4	5	6	7	8	9	10	11	12
Character (n-z)	n	o	p	q	r	s	t	u	v	w	x	y	z
Value of h$_{char}$	13	14	15	16	17	18	19	20	21	22	23	24	25

Table 1: Values of h_{char} for all the English characters.

We can then use this function to generate another simple hash function, this time for strings:

$$h_{str}(s) = h_{char}(s[0]) \bmod M$$

This hash function is *not very good*, particularly when compared to the default implementation of `String.hashCode()` in `Java`. First of all, every pair of lowercase strings which begins with the same letter will collide. This is true even if $M > 26$, the cardinality of the English alphabet! But it's not of course just the *immediate* collisions that cause us grief: the first character collisions tend to make "clusters" in the table which make even insertions for strings that begin with a **new** first character (when compared to the first characters of the already inserted strings) collide! See Figure 6 for an example of this. Note that the hash table is reasonably large that no re-sizing is necessary during the insertions we show ... even *barely* so.

0	rocket	9	
1	sun	10	
2	sight	11	
3	torus	12	
4	elated	13	
5	feather	14	
6	fiscal	15	
7	fang	16	
8	gorilla		

Figure 6: An example of the "clustering" phenomenon in LP when using a low quality hash function. Even keys with novel first characters collide with existing keys, because the hash function "targets" already congested areas of the table. Sequence of insertions: *sun, elated, sight, rocket, torus, feather, fiscal, fang, gorilla.*

In Figure 6 we see that inserting several keys with the same first character (*sun, sight, feather, fiscal, fang*) **enlarges the relevant collision chain**. But it's not just their *own* collision chain that they enlarge, but also **that of other keys** (*torus, gorilla*), which do **not** hash to the same bucket!

Unfortunately, with the simple collision resolution technique that LP employs, we **cannot** hope to alleviate the clustering phenomenon. Our only solution to it is **re-sizing the table when we have to**. Do note, however, that with a hash function this bad, **even re-sizing cannot help us**, since the operation *mod M* does **not** change the problems of h_{str}...

4.3   Ordered Linear Probing

But there **is** something we can do to improve our fortune. Consider for a moment a search for "entropy" in the table of Figure 6. Since we are humans and can immediately see the entirety of the table, we know that this search is *destined to fail*. But *how many probes* will the insertion algorithm need to determine this? It would need **six** probes (the final one hitting `null` in cell 9), and this is because of one collision with a fellow word that begins with an 'e', followed by the existence of keys with a first character of 'f', which did not even hash to the same position as us any way! Not to mention that the "f's pushed **gorilla** over and that made things even worse for us. If only we could make those keys "get out of the way" so we can *fail this search faster* and move on to operations not destined to fail!

A simple modification of Linear Probing, **Ordered** Linear Probing (**OLP** for short) tries to achieve exactly this goal. It alleviates this problem by keeping the collision chains *ordered*. The way that this method works is as follows: consider that at some point in the collision resolution process that LP employs, the new key k' encounters a key k which has the property that $k < k'$ (**strictly** smaller). The comparison operator $<$ here is assumed to mean whatever "$<$" means for the provided key type: numerical comparison, alphabetical comparison, custom `compareTo()`, etc. Then, the algorithm will put k' in the position of k and continue the insertion *as if the key to be inserted is k!* Effectively, the insertion

11

algorithm will keep going down the collision chain for an empty spot to put k in. We are thus keeping the collision chain *ordered* **and** we do **not** break the search for k, since we still have a *contiguous cluster of collision chains* which allows the algorithm to **not** reach `null` before it finds k.

Figure 7 shows an example of what the hash table depicted in Figure 6 would look like if we had employed OLP instead of simple LP. We once again see that we have **not** alleviated the problem of **clustering**; what we *have* done is make searches *destined* to fail, fail **faster**!

0	rocket	9	
1	sight	10	
2	sun	11	
3	torus	12	
4	elated	13	
5	fang	14	
6	feather	15	
7	fiscal	16	
8	gorilla		

Figure 7: The result of inserting the numbers in the sequence of Figure 4 but this time employing *Ordered Linear Probing* as our collision resolution technique.

You should implement this method in the class `OrderedLinearProbingHashTable`. You might find that several of the methods you implement are 100% identical to those employed by `LinearProbingHashTable`. If that is the case, we would recommend that you refactor your code such that methods with identical definitions are **all** merged into **one** `protected` method in `OpenAddressingHashTable`, so you only have to debug that **one** method if something bad were to happen.

4.4 Quadratic Probing

In lecture we saw that Linear Probing is susceptible to the "clustering" phenomenon, where various different collision chains end up "crowding" next to each other and even "overlapping". This causes several collisions for **even wildly** different hash codes when compared to the ones that started the chains. We also saw that tuning Linear Probing such that its "jump" is changed from 1 to some other number, e.g 2 or 3, does **not** solve the clustering problem: instead, the clusters become *discontiguous* on the table.

This begs the question: *what if, instead of having a static offset to Linear Probing, we were to* **increase** *the "step" that the algorithm takes* **every time it encounters a collision?** One studied solution that implements this idea is **quadratic probing (QP)**. QP, in its simplest form (which is the one you will implement in this project), employs the following memory allocation function m_{qp}:

$$m_{qp}(k, i) = \big(h(k) + (i - 1) + (i - 1)^2 \big) \bmod M$$

12

which will lead into the following memory addresses being probed:

$$h(k) \bmod M, (h(k) + 2)) \bmod M, (h(k) + 6) \bmod M, (h(k) + 12) \bmod M, \ldots$$

Note that the offset is **always** computed from the address that $h(k)$ initially probed. For example, if the table of Figure 5 was using **quadratic** instead of linear probing, we would have the single memory allocations shown in Figure 8.


Figure 8: Example of insertion speed-up achieved by quadratic probing; $m_{qp}(22, 2)$ attains a bigger "jump" than $m_{lp}(22, 2)$ and reduces the number of probes required to insert 22 by 1. For practice, try inserting 11 into either one of the hash tables **after** you insert 22, and look at the speed-up attained by QP! ☺

The entire idea behind QP is that if a key collides with another, we need to *try harder to make it not collide in the immediate future.* By increasing the quadratic "step" every time that a collision happens, the algorithm hopes to disperse keys that collide more **aggressively**. Feel free to read the relevant Wikipedia article or scour the web for additional resources on how QP improves upon LP. It doesn't improve **universally**, though. For one, it doesn't display the **cache locality** that Linear Probing displays, especially for keys that collide a lot (so the value of i is large for them). Also, the hash function is just a tiny bit more **expensive** to compute, since there's another summand and a squaring involved. Perhaps a form of *caching* could be employed to make it cheaper to compute.

Two final implementation notes in QP. First, you might once again find that a lot of the code you write is identical to that of the other classes (though arguably less so). We would once again encourage you to package common pieces of code as `protected` methods in `OpenAddressingHashTable`. Second, since it is **very hard** to define the contiguous clusters when the key's "jump" changes with every collision encountered, in hard deletions you should simply **reinsert all keys besides the key that you want to delete**. Sounds inefficient, **is** inefficient. For this reason...

5 Soft vs Hard deletion in Openly Addressed Hash Tables

As you know from lecture, during hard deletion, LP nullifies and re-inserts all keys that follow it *within the cluster itself* (so it should stop the process when it hits `null`). The same process is followed by OLP, whereas QP takes it one step "further" by simply re-inserting **all** keys that are **not** the key of interest, since in QP it is not straightforward to determine the memory topology of the clusters. On the other hand, this is **not** a problem that affects `SeparatelyChainedHashTable`, since those kinds of tables can simply call `KVPairList.remove()` and be done with a key with reasonable efficiency and no need to re-insert any other keys!

For this reason, we introduce a named constant called `TOMBSTONE` in `OpenlyAddressed HashTable`. This constant will be used as a placeholder for *softly deleted* `KVPair`s in the table. It is **supremely important** to understand the following things about this constant:

- Any memory address that holds this constant is an *available* position for insertion of a new key. Therefore, for purposes of **insertion**, if the hash table in question soft-deletes, you can treat a tombstone-containing cell as an empty cell.

- Any memory address that holds this constant is **not** a `null` cell! So, the collision chains that "went over" this memory address before the soft deletions still continue to "go over it"! We **do not break** collision chains this way(this would be bad)! In fact, since we know that just `null`-ifying collision chains breaks search, yielding the necessity of re-insertion of the subsequent cluster elements, we are **obligated** to **not** treat `TOMBSTONE` as a `null` entry. You should take this into consideration when implementing `get()` in your hash tables.

- Since tombstone-containing memory addresses will still burden searches with an additional probe, **they are considered as an occupied cell when we resize**. For example, in Figure 9, any new insertion will contribute to a resizing, since the hash table **already** has a count of elements (`size()`) of **more than half** its capacity (`capacity()`). However, **TOMBSTONES THEMSELVES DO NOT GET RE-INSERTED!** What kind of hash value could we expect off of a tombstone to re-insert it, anyway? [2]

---

[2]It turns out that tombstones are quite the popular idea in Computer Science as a whole and in hashing in particular.

Left table:

0	walrus
1	lion
2	RIP
3	lynx
4	dodo
5	RIP
6	
7	**hippo**
8	
9	
10	

$e(11) = 19$

Right table:

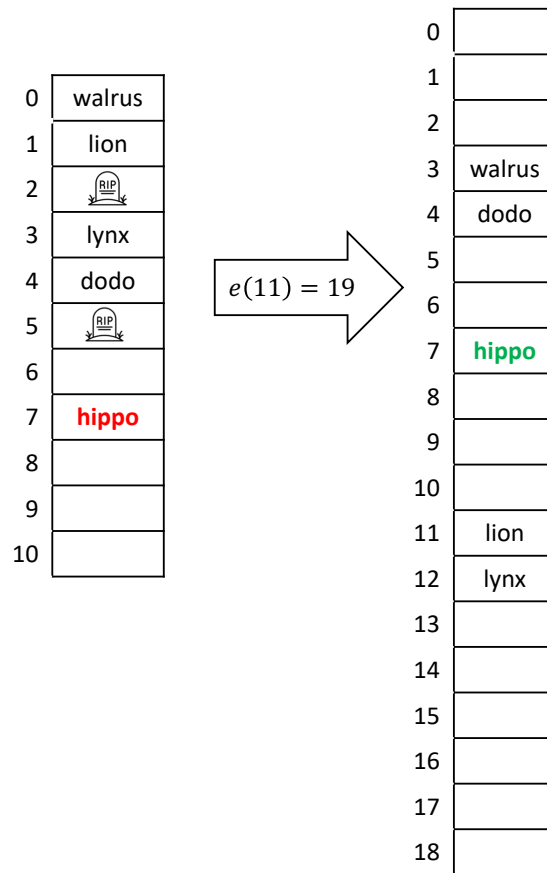| 0 | |
1	
2	
3	walrus
4	dodo
5	
6	
7	**hippo**
8	
9	
10	
11	lion
12	lynx
13	
14	
15	
16	
17	
18	

Figure 9: Illustration of the fact that "tombstones" in a hash table contribute to its count of elements.

When **hippo** is inserted, the table *might* contain only two *actual* keys, but the tombstones are actually *there*, so the resizing operation will be triggered *before* **hippo** is inserted. Why can't the tombstones just be "swept away" by essentially being treated like `null` entries? Well, because that would break search for subsequent keys! In this case, if the "dummy" tombstone entry at position 2 were treated like a `null` entry, future searches for **lynx** would **falsely fail!**

Finally, how do we tune an `OpenAddressingHashTable` instance for hard or soft deletion? Simple. The various constructors of the classes that extend this class have a `boolean` argument which, if `true`, determines that the instance being created will perform **soft** deletion, **hard** otherwise. You can even run your own timing experiments this way, and compare hard vs soft deletion's efficiency as well as how much they affect the efficiency of *other* operations!

6   FAQs

Q: *Why is it that* `SeparateChainingHashTable` *has two* `public` *methods (*`enlarge()`*,* `shrink()`*) which are* **not** *part of the interface* `HashTable`*?*

A: Because enlarging or reducing the number of entries in a hash table implemented with Separate Chaining as its collision resolution strategy is a process that never *has* to happen **automatically** in order for its operations to work (particularly, insertions). Enlarging the hash table can lead to better *efficiency* of operations, while reducing its size can lead to better storage tradeoffs after numerous **deletions** have happened. This means that changing the Separately Chained hash table's *capacity* is an issue that **should be left with the caller to decide**. Maybe the caller decides to enlarge when the capacity is at 70%; if so, the caller must **explicitly** make the call to `enlarge()` (similarly for `shrink()`). On the other hand, an openly addressed hash table will need to **internally resize** the table in order to **not just allow for better performance and storage trade-offs** but, in the case of insertions, **to even allow for the operation to complete**!

Q: *Does this mean that I* **don't** *need such methods for my Openly Addressed Hash Tables, that is,* `LinearProbingHashTable` *and* `QuadraticProbingHashTable`*?*

A: You will **absolutely** need such methods, but they have no business being `public` methods.

Q: *How should resizings be implemented?*

A: To begin with, you will **not** need to resize during **deletions**. Theferore, for insertions, you should follow the approach that we have discussed in class: **the <u>first</u> insertion that takes place <u>after</u> your hash table is at 50% capacity or more should <u>first</u> trigger a resizing of the hash table to the <u>largest</u> prime number <u>smaller</u> than <u>twice your current size</u> (**`PrimeGenerator.getNextPrime()` **takes care of this for you), and <u>then</u> insert the new key.** For example, if the current hash table capacity is 7, the 4th insertion will cause the hash table to have a count of 4, which is $\approx 57.1\%$ of the hash table size. We will **not** resize after the 4th insertion though, because we **don't know whether a 5th one will come yet**, and it is possible that we would be resizing "for nothing". **If** a 5th insertion is requested, we will **first** resize to 13, the largest prime number smaller than $2 \cdot 7 = 14$, and **then** we will insert the 5th element. Note that the element might be hashed to an address different than the one it would have been hashed to if we had **not** resized, since the hash code will be "modded" by a **new** hash table size. It is important that you stick **exactly** to this guideline, because `HashTable` instances expose a `public` method called `capacity()` which checks for the actual hash table **size**, i.e the number of cells of the internal 1D array that implements the actual table **whether they are**
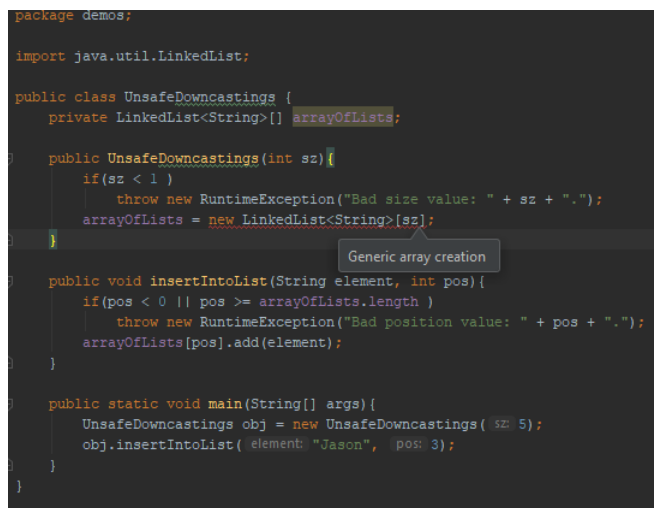
occupied or not. This means that we can **test** for the return value of that method, and we will be *expecting* that you follow these guidelines to a tee. Feel free to use the class `utils.PrimeGenerator` to get the appropriate prime numbers for free. Check the file `StudentTests.java` for some examples of how we can test for the return value of `capacity()`.

Q: *For an Openly Addressed Hash Table that is **very sparse**, doesn't it make sense to truncate its size to the first prime greater than its current number of elements, such that we save space?*

A: Yes, in practice, you should. However, in this project, we assume that our hash tables are kept at a reasonable load factor so that you don't ever encounter significant sparsity.

Q: *Why did you implement your own linked list over `KVPair` instances (`KVPairList`) instead of just instantiating the `private` data field `table` of `SeparateChainingHashTable` with a `java.util.LinkedList<KVPair>`? Surely that is easier to do instead of writing your own list for the project and then testing it!*

A: We do this because *creating a raw array over **generic** types* in Java is a *pain*. As you can see in figure 10, in Java it is **not** possible to create a raw array over generics. `KVPairList` is, unfortunately, a generic type since it extends `Iterable<KVPair>`, itself a generic. We have included two example source code `Java` files, `GenericArrays.java` and `UnsafeDowncastings.java`, under `src/demos` for you to consult, run, and understand this problem better.

```java
package demos;

import java.util.LinkedList;

public class UnsafeDowncastings {
    private LinkedList<String>[] arrayOfLists;

    public UnsafeDowncastings(int sz){
        if(sz < 1 )
            throw new RuntimeException("Bad size value: " + sz + ".");
        arrayOfLists = new LinkedList<String>[sz];
    }
                                    Generic array creation
    public void insertIntoList(String element, int pos){
        if(pos < 0 || pos >= arrayOfLists.length )
            throw new RuntimeException("Bad position value: " + pos + ".");
        arrayOfLists[pos].add(element);
    }

    public static void main(String[] args){
        UnsafeDowncastings obj = new UnsafeDowncastings( sz: 5);
        obj.insertIntoList( element: "Jason", pos: 3);
    }
}
```

Figure 10: Creating an array of generic types leads to a **compile-time** error.

Q: *What courses should I take after CMSC420?*

A: You should take the courses you feel that you would *like*, as long as you can find seats in them. That said, if you are interested in courses that will mesh well with what you

have learned in 420 and perhaps even offer you alternative viewpoints, consider **451** and **424**.

Q: *What places should I visit in Greece if I ever travel there?*

A: In reality, everywhere you go will be a fantastic experience, especially if you haven't been in Europe before. Mykonos and Santorini (Thera, the one with the volcano) are tourist traps, but they are gorgeous tourist traps (the latter more so than the party-capital former) and they offer such a great experience to somebody who hasn't seen this stuff before that I cannot tell you to avoid them merely because they are tourist traps...

If you want some reliable and quality tourism infrastructure and simultaneously want to avoid tourist traps, I would recommend spending a few days in **Athens**, a millenia-old city with a **ton** of historical tours, museums and a **ginormous** selection of **terrific** food places, views and regular big city hustle and bustle (6 million people). I would then select an island and **avoid** a hotel/motel/AirBnb in a tourist trap area. For example, if you choose to go to Rhodes, don't stay in Faliraki. If you choose to go to Crete, avoid St. Nicholas and Vai. If you go to Kefalonia, avoid Makris Gialos. In Zakynthos, avoid Laganas. Ask your guide for more current information. Do **not** waste your time on Thessalonica, Patras, Larissa, Lamia or Sparta. Those places have **nothing** for a foreign person to enjoy. Don't tell them I told you, because non-Athenian people hate Athenians (I am an Athenian) yet simultaneously everybody flocks to Athens because the only industries outside Athens are tourism, agriculture and food.

Figure 11: Hades and Zeus are slowly, yet deliberately and unflinchingly, judging you based on your performance in this project.